

SASSA SQL Workshop

Fall 2019

Why DBMS

	A	B
1	Country	City
2	Canada	Toronto
3	Canada	Vancouver
4	Canada	Montréal
5	Canada	Ottawa
6	Canada	Edmonton
7	Canada	Victoria
8	Canada	Halifax
9	Canada	St. John's
10	Canada	Québec
11	USA	Seattle
12	USA	New York
13	USA	Portland
14	USA	Denver
15	USA	Austin
16	USA	San Francisco

Example

Search for

County == "USA" && City ==
"Austin"

Average case: $O(n)$

INDEXING

	A	B
1	country_id	Country
2		1 Canada
3		2 USA

Search for

County == "USA" && City == "Austin"

Average case: $O(\log(n))$

	A	B	C
1	city_id	City	country_id
2		1 Toronto	1
3		2 Vancouver	1
4		3 Montréal	1
5		4 Ottawa	1
6		5 Edmonton	1
7		6 Victoria	1
8		7 Halifax	1
9		8 St.John's	1
10		9 Québec	1
11		10 Seattle	2
12		12 New York	2
13		13 Portland	2
14		14 Denver	2
15		15 Austin	2
16		16 San Francisco	2

SCHEDULING

Example

Two people buying MOVIE TICKETS

What if the data was stored in a CSV...

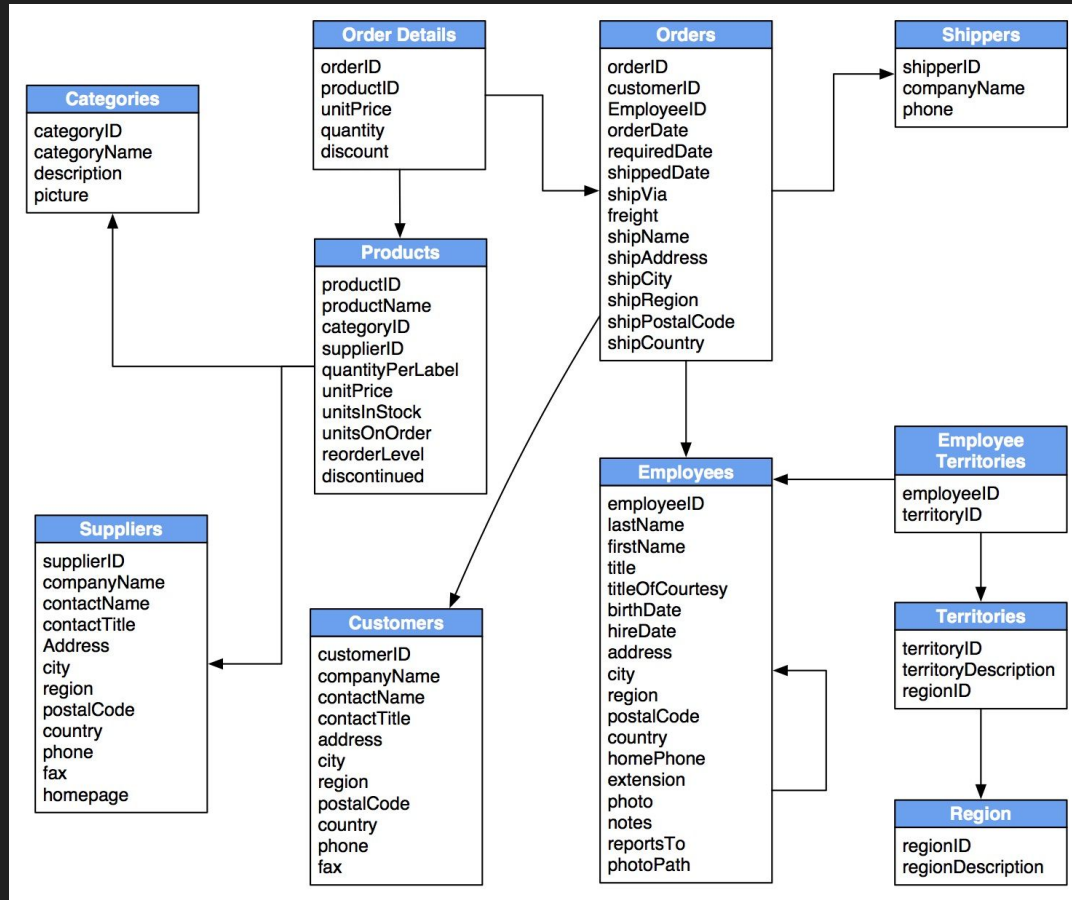
Can we WRITE to this document simultaneously?

Whose data is added first?


RELATIONS

ER diagrams

dbdiagram.io



ENTITY RELATIONSHIP DIAGRAM



OUR ER DIAGRAM

We will have 7 **ENTITIES** :

FILM - stores films data such as title, release year, length, rating, etc.

CATEGORY - stores film's categories data.

FILM_CATEGORY - stores the relationships between films and categories.

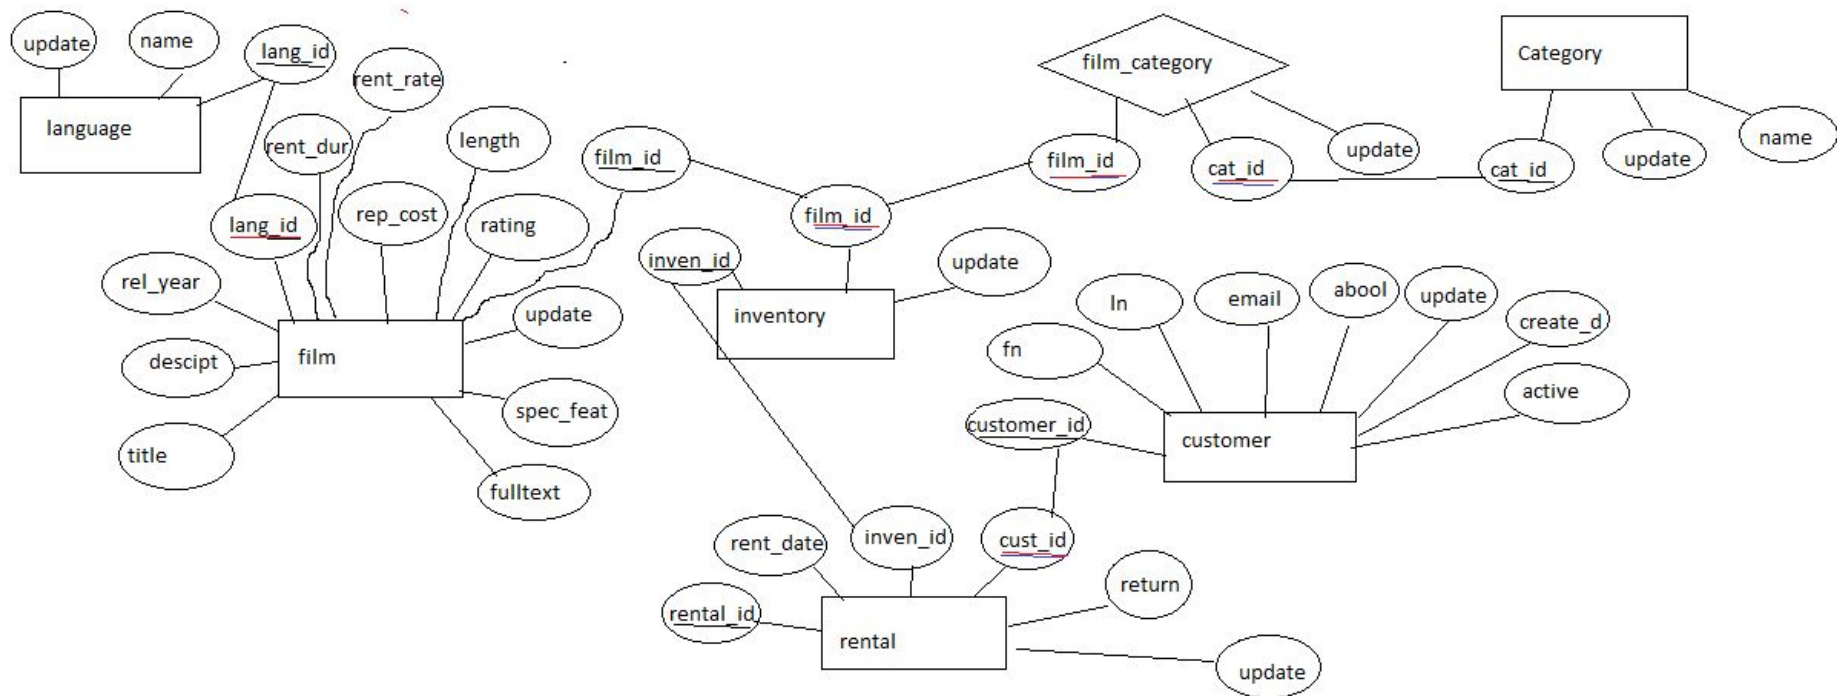
INVENTORY - stores inventory data.

RENTAL - stores rental data.

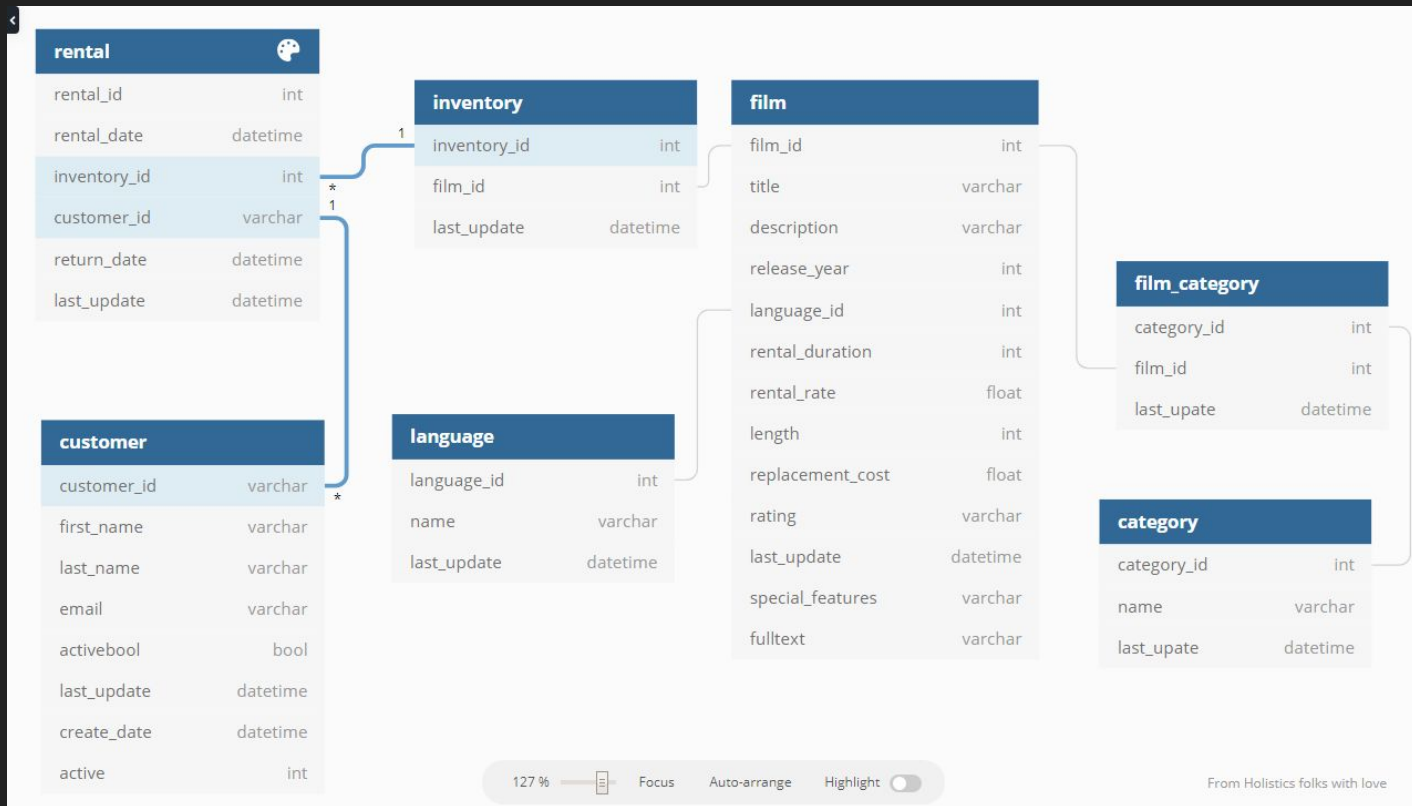
CUSTOMER - stores customers data.

LANGUAGE - stores the film language

OUR ER DIAGRAM

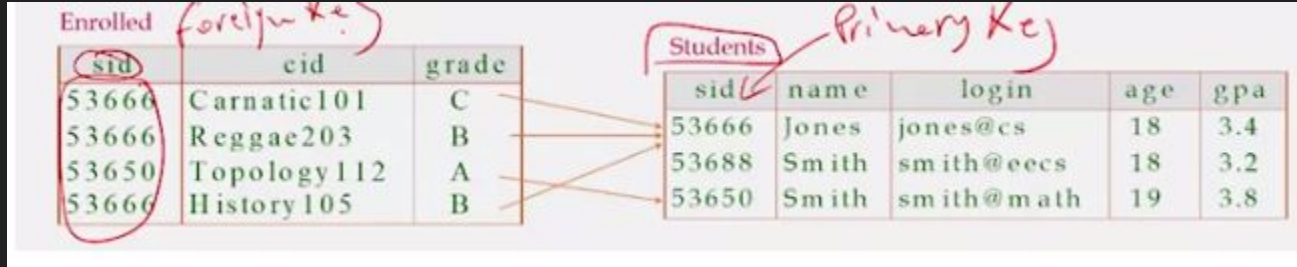


OUR RELATIONSHIP MODEL



KEYS

EXAMPLE



Self study KEY CONSTRAINTS

CREATE TABLE

```
CREATE TABLE rental (  
    rental_id SERIAL,  
    rental_date timestamp with time zone NOT NULL,  
    inventory_id integer NOT NULL,  
    customer_id integer NOT NULL,  
    return_date timestamp with time zone ,  
    last_update timestamp with time zone DEFAULT now() NOT NULL,  
    PRIMARY KEY (rental_id),  
    FOREIGN KEY (customer_id) REFERENCES customer  
);
```

INSERT

```
INSERT INTO inventory VALUES (612, 133, '2017-02-15 10:09:17');
```

SELECT Statements

SELECT statements are used to select data from a database

SELECT

-- All columns




*

FROM

-- The Schema 'public'

-- The Table 'country'

public.country

Data Output		Explain	Messages	Notifications
	 country_id [PK] integer 	country character varying (50) 	last_update timestamp without time zone 	
1	1	Afghanistan	2006-02-15 09:44:00	
2	2	Algeria	2006-02-15 09:44:00	
3	3	American Samoa	2006-02-15 09:44:00	
4	4	Angola	2006-02-15 09:44:00	
5	5	Anguilla	2006-02-15 09:44:00	
6	6	Argentina	2006-02-15 09:44:00	
7	7	Armenia	2006-02-15 09:44:00	
8	8	Australia	2006-02-15 09:44:00	
9	9	Austria	2006-02-15 09:44:00	
10	10	Azerbaijan	2006-02-15 09:44:00	

SELECT Statements

`SELECT` statements are used to select data from a database

```
SELECT
```

```
-- The column country
```


```
country
```

```
FROM
```

```
-- The Schema 'public'
```

```
-- The Table 'country'
```

```
public.country
```

Data Output		Explain	Messages	Notifications
	 country character varying (50)			
1	Afghanistan			
2	Algeria			
3	American Samoa			
4	Angola			
5	Anguilla			
6	Argentina			
7	Armenia			
8	Australia			
9	Austria			
10	Azerbaijan			

SELECT Statements

SELECT

*

FROM

public.country,
public.city

Two tables are included in this SELECT statement

Does anyone see what's wrong?

Data Output	Explain	Messages	Notifications				
	country_id integer	country character varying (50)	last_update timestamp without time zone	city_id integer	city character varying (50)	country_id smallint	last_update timestamp without time zone
1	1	Afghanistan	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
2	2	Algeria	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
3	3	American Samoa	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
4	4	Angola	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
5	5	Anguilla	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
6	6	Argentina	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
7	7	Armenia	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
8	8	Australia	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
9	9	Austria	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
10	10	Azerbaijan	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25

SELECT Statements

SELECT

*

FROM

public.country,
public.city

Two tables are included in this SELECT statement

Does anyone see what's wrong?

Data Output	Explain	Messages	Notifications				
	country_id integer	country character varying (50)	last_update timestamp without time zone	city_id integer	city character varying (50)	country_id smallint	last_update timestamp without time zone
1	1	Afghanistan	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
2	2	Algeria	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
3	3	American Samoa	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
4	4	Angola	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
5	5	Anguilla	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
6	6	Argentina	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
7	7	Armenia	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
8	8	Australia	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
9	9	Austria	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
10	10	Azerbaijan	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25

SELECT Statements

SELECT

*

FROM

public.country,

public.city

WHERE

country.country_id = city.country_id

Data Output Explain Messages Notifications

	country_id integer	country character varying (50)	last_update timestamp without time zone	city_id integer	city character varying (50)	country_id smallint	last_update timestamp without time zone
1	87	Spain	2006-02-15 09:44:00	1	A Corua (La Corua)	87	2006-02-15 09:45:25
2	82	Saudi Arabia	2006-02-15 09:44:00	2	Abha	82	2006-02-15 09:45:25
3	101	United Arab Emirates	2006-02-15 09:44:00	3	Abu Dhabi	101	2006-02-15 09:45:25
4	60	Mexico	2006-02-15 09:44:00	4	Acua	60	2006-02-15 09:45:25
5	97	Turkey	2006-02-15 09:44:00	5	Adana	97	2006-02-15 09:45:25

SELECT Statements

SELECT




country.country,
city.city

FROM

public.country,
public.city

WHERE

country.country_id = city.country_id

	Data Output	Explain	Messages	Notifications
	 country character varying (50)			city character varying (50) 
1	Spain			A Corua (La Corua)
2	Saudi Arabia			Abha
3	United Arab Emirates			Abu Dhabi
4	Mexico			Acua
5	Turkey			Adana
6	Ethiopia			Addis Abeba
7	Yemen			Aden
8	India			Adoni
9	India			Ahmadnagar
10	Japan			Akishima

SELECT Statements - Equi Join

These aren't very useful results though.

Let's say we wanted to find out all the people who have rented the movie 'Caddyshack Jedi'

SELECT Statements - Equi Join

```
SELECT
```

```
*
```

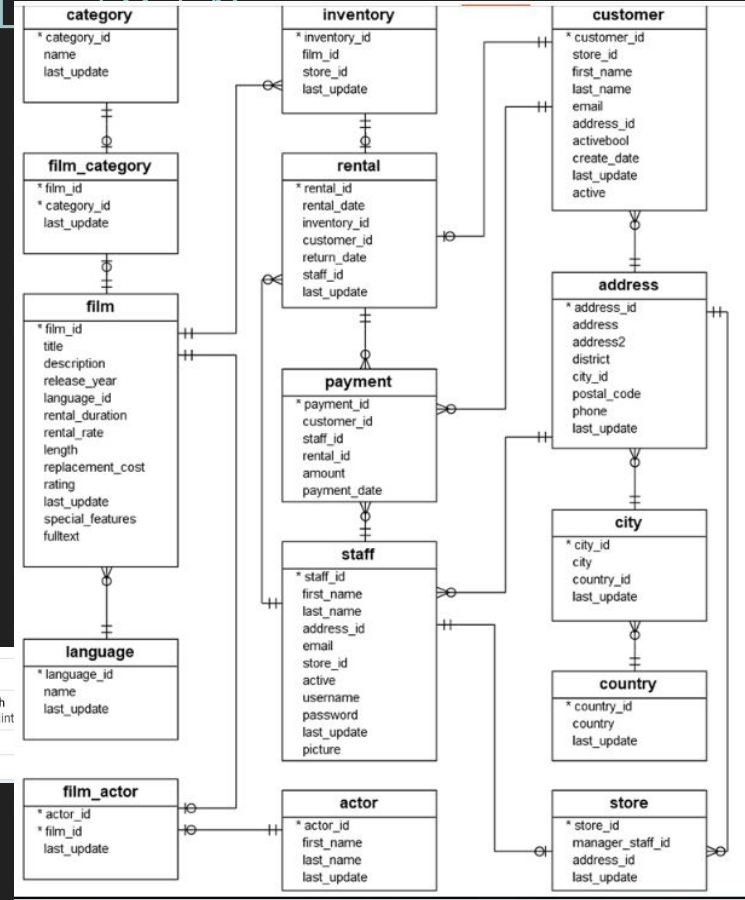
```
FROM
```

```
public.film
```

```
WHERE
```

```
public.film.title = 'Caddyshack Jedi'
```

Data Output								Explain	Messages	Notifications					
	film_id [PK] integer		title character varying (255)		description text		release_year integer		language_id smallint		rental_duration smallint		rental_rate numeric (4,2)		length smallint
1		111	Caddyshack Jedi		A Awe-Inspiring Epistle of a Woman And a Madman ...		2006				1		3		0.99



SELECT Statements - Equi Join

SELECT

```
public.film.title
```

FROM

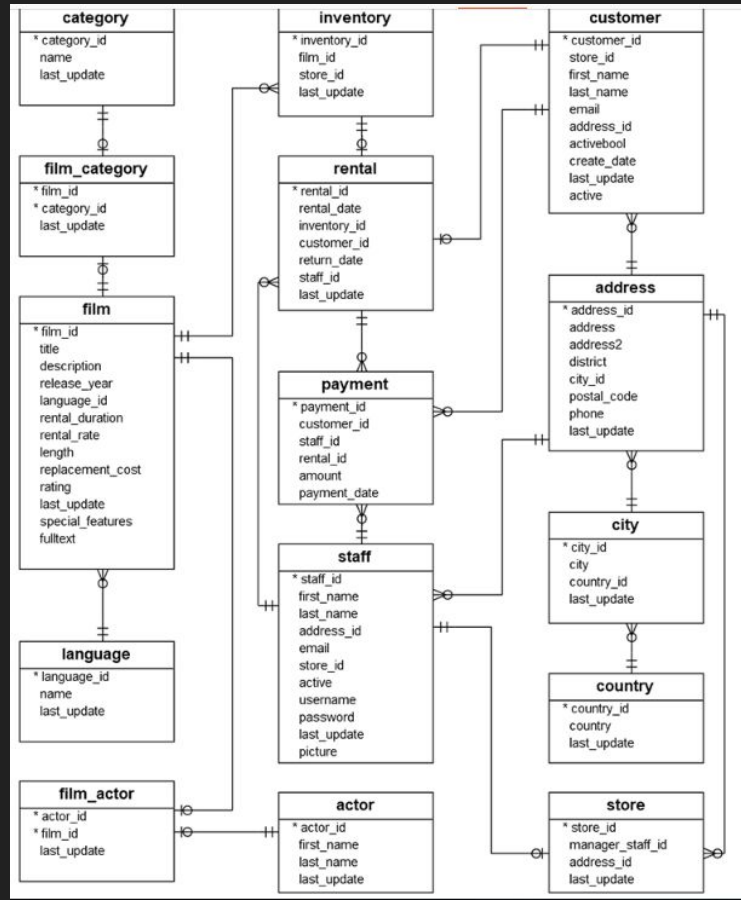
```
public.film,
```

```
public.inventory
```

WHERE

```
public.film.film_id = public.inventory.film_id
```

```
AND public.film.title = 'Caddyshack Jedi'
```



SELECT Statements - Equi Join

```
SELECT
    public.film.title
FROM
    public.film,
    public.inventory,
    public.rental
WHERE
    public.film.film_id = public.inventory.film_id
    AND public.inventory.inventory_id = public.rental.inventory_id
    AND public.film.title = 'Caddyshack Jedi'
```

SELECT Statements - Equi Join

```
SELECT
    public.film.title
FROM
    public.film,
    public.inventory,
    public.rental,
    public.customer
WHERE
    public.film.film_id = public.inventory.film_id
    AND public.inventory.inventory_id = public.rental.inventory_id
    AND public.rental.customer_id = public.customer.customer_id
    AND public.film.title = 'Caddyshack Jedi'
```

SELECT Statements - Equi Join

```
SELECT
    public.film.title
FROM
    public.film
    public.inventory
    public.rental
    public.customer
WHERE
    public.film.film_id = public.inventory.film_id
    AND public.inventory.inventory_id = public.rental.inventory_id
    AND public.rental.customer_id = public.customer.customer_id
    AND public.film.title = 'Caddyshack Jedi'
```

Data Output		Explain	Messages	Notifications
	title character varying (255)			
1	Caddyshack Jedi			
2	Caddyshack Jedi			
3	Caddyshack Jedi			
4	Caddyshack Jedi			
5	Caddyshack Jedi			
6	Caddyshack Jedi			
7	Caddyshack Jedi			
8	Caddyshack Jedi			
9	Caddyshack Jedi			
10	Caddyshack Jedi			
11	Caddyshack Jedi			
12	Caddyshack Jedi			
13	Caddyshack Jedi			
14	Caddyshack Jedi			
15	Caddyshack Jedi			
16	Caddyshack Jedi			

SELECT Statement

SELECT

```
public.film.title,  
public.customer.first_name,  
public.customer.last_name
```

FROM

```
public.film,  
public.inventory,  
public.rental,  
public.customer
```

WHERE

```
public.film.film_id = public.inventory.inventory_id  
AND public.inventory.inventory_id = public.rental.inventory_id  
AND public.rental.customer_id = public.customer.customer_id  
AND public.film.title = 'Caddyshack Jedi'
```


Data Output Explain Messages Notifications

	title character varying (255)	first_name character varying (45)	last_name character varying (45)
1	Caddyshack Jedi	Sheila	Wells
2	Caddyshack Jedi	Tom	Milner
3	Caddyshack Jedi	George	Linton
4	Caddyshack Jedi	Lydia	Burke
5	Caddyshack Jedi	Angel	Barclay
6	Caddyshack Jedi	Charlie	Bess
7	Caddyshack Jedi	Neil	Renner
8	Caddyshack Jedi	Louis	Leone
9	Caddyshack Jedi	George	Linton
10	Caddyshack Jedi	Dianne	Shelton
11	Caddyshack Jedi	Charlene	Alvarez
12	Caddyshack Jedi	Edith	Mcdonald
13	Caddyshack Jedi	Charlene	Alvarez
14	Caddyshack Jedi	Jimmy	Schrader
15	Caddyshack Jedi	George	Linton
16	Caddyshack Jedi	Dawn	Sullivan

SELECT Statements

There are several types of joins for SELECT statements

SQL EQUI JOIN
SQL NON EQUI JOIN
SQL INNER JOIN
SQL NATURAL JOIN
SQL CROSS JOIN
SQL OUTER JOIN
SQL LEFT JOIN
SQL RIGHT JOIN
SQL FULL OUTER JOIN
Join a table to itself
SQL SELF JOIN



SELECT Statements

In the interest of time, I'll only demonstrate a LEFT JOIN

SELECT Statements - Left Join

```
SELECT
    public.customer.first_name,
    public.customer.last_name,
    public.rental.rental_id,
    public.customer.customer_id
FROM
    public.customer
LEFT JOIN public.rental
    ON public.customer.customer_id = public.rental.customer_id
```

SELECT Statements - Left Join

SELECT

```
public.customer.first_name,  
public.customer.last_name,  
public.rental.rental_id,  
public.customer.customer_id
```

FROM

```
public.customer
```

```
LEFT JOIN public.rental
```

```
ON public.customer.customer_id = public.rental.customer_id
```

WHERE

```
rental_id IS NULL
```

Data Output	Explain	Messages	Notifications				
<div><div>▼</div><div>first_name character varying (45)</div></div>	<div><div>🔒</div></div>	<div><div>last_name character varying (45)</div></div>	<div><div>🔒</div></div>	<div><div>rental_id integer</div></div>	<div><div>🔒</div></div>	<div><div>customer_id integer</div></div>	<div><div>🔒</div></div>

SELECT Statements - Left Join

```
INSERT INTO
```

```
  Customer
```

```
VALUES (
```

```
  600,
```

```
  1,
```

```
  'Lily',
```

```
  'Potter',
```

```
  'lily_potter@hogwarts.com',
```

```
  1,
```

```
  true,
```

```
  now(),
```

```
  now()
```

```
)
```

Data Output

Explain

Messages

Notifications

INSERT 0 1

Query returned successfully in 44 msec.

SELECT Statements - Left Join

```
SELECT
    public.customer.first_name,
    public.customer.last_name,
    public.rental.rental_id,
    public.customer.customer_id
FROM
    public.customer
LEFT JOIN public.rental
    ON public.customer.customer_id = public.rental.customer_id
WHERE
    rental_id IS NULL
```

	Data Output	Explain	Messages	Notifications
	first_name character varying (45)	last_name character varying (45)	rental_id integer	customer_id integer
1	Lily	Potter	[null]	600

SELECT Statements - Left Join

SELECT

```
public.customer.first_name,  
public.customer.last_name,  
public.rental.*
```

FROM

```
public.customer
```

```
LEFT JOIN public.rental
```

```
ON public.customer.customer_id = public.rental.customer_id
```

WHERE

Data Output	Explain	Messages	Notifications										
	first_name character varying (45)	last_name character varying (45)	rental_id integer	rental_date timestamp without time zone	inventory_id integer	customer_id smallint	return_date timestamp without time zone	staff_id smallint	last_update timestamp without time zone				
1	Lily	Potter	[null]	[null]	[null]	[null]	[null]	[null]	[null]				

SELECT Statements - Left Join

SELECT

public Lily has never rented a movie, hence all of her
public results are expectedly NULL
public

FROM

public

LEFT JOIN

ON p

WHERE

Data Output	Explain	Messages	Notifications								
	first_name character varying (45)	last_name character varying (45)	rental_id integer	rental_date timestamp without time zone	inventory_id integer	customer_id smallint	return_date timestamp without time zone	staff_id smallint	last_update timestamp without time zone		
1	Lily	Potter	[null]	[null]	[null]	[null]	[null]	[null]	[null]		

SELECT Statements - Some For You To Try

See if you can come up with SELECT statements for the following questions:

1. Find the First and Last Names of the customers who have rented 'Caddyshack Jedi' OR 'Born Spinal'

SELECT Statements - Some For You To Try

See if you can come up with SELECT statements for the following questions:

2. Find the name of all the customers from 'Vancouver'

SELECT Statements - Some For You To Try

See if you can come up with SELECT statements for the following questions:

3. Find all the customers who have paid more than \$10